

```
In [1]: from utils.data import Dataset
from utils.conllu import read_conllu_dataset
import numpy as np

In [2]: dataset = Dataset()

train_sentences = read_conllu_dataset("data/ro_rrt-ud-train.conllu")
test_sentences = read_conllu_dataset("data/ro_rrt-ud-test.conllu")

# Fit on training data
X_train, y_train = dataset.fit(train_sentences, mode="chars")

# Encode test data (fixed shape)
X_test, y_test = dataset.encode(test_sentences)

In [3]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
Out[3]: ((8043, 163, 55), (8043, 163, 17), (729, 163, 55), (729, 163, 17))

In [4]: import tensorflow as tf

In [5]: _, seq_len, char_feat_len = X_train.shape
output_dim = y_train.shape[-1]

In [6]: sample_weight = np.where(np.argmax(y_train, axis=-1) != 0, 1.0, 0.0)

In [7]: input_layer = tf.keras.layers.Input(shape=(seq_len, char_feat_len))
lstm = tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(128, return_sequences=True))(input_layer)
output_layer = tf.keras.layers.TimeDistributed(tf.keras.layers.Dense(output_dim, activation='softmax'))(lstm)

model = tf.keras.models.Model(inputs=input_layer, outputs=output_layer)
model.compile(optimizer=tf.keras.optimizers.Adam(0.001), loss='categorical_crossentropy', weighted_metrics=['accuracy'])

model.summary()
```

Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 163, 55)	0
bidirectional (Bidirectional)	(None, 163, 256)	188,416
time_distributed (TimeDistributed)	(None, 163, 17)	4,369

Total params: 192,785 (753.07 KB)

Trainable params: 192,785 (753.07 KB)

Non-trainable params: 0 (0.00 B)

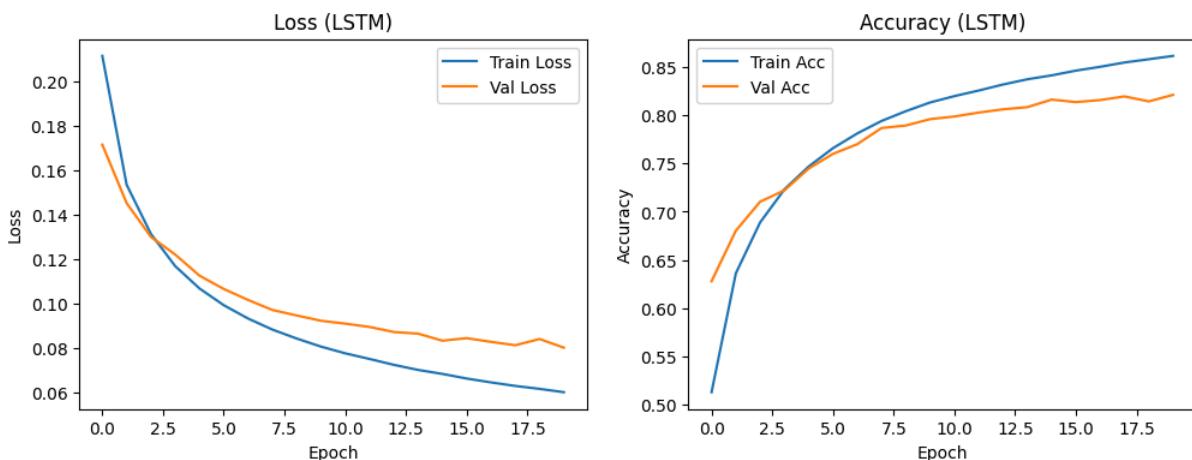
```
In [8]: history = model.fit(
    X_train, y_train,
    validation_split=0.1,
    epochs=20,
    batch_size=64,
    sample_weight = sample_weight
)

Epoch 1/20
114/114 153s 1s/step - accuracy: 0.4310 - loss: 0.2563 - val_accuracy: 0.6280 - val_loss: 0.1717
Epoch 2/20
114/114 129s 1s/step - accuracy: 0.6226 - loss: 0.1600 - val_accuracy: 0.6803 - val_loss: 0.1453
Epoch 3/20
114/114 130s 1s/step - accuracy: 0.6784 - loss: 0.1348 - val_accuracy: 0.7104 - val_loss: 0.1301
Epoch 4/20
114/114 131s 1s/step - accuracy: 0.7195 - loss: 0.1189 - val_accuracy: 0.7222 - val_loss: 0.1221
Epoch 5/20
114/114 130s 1s/step - accuracy: 0.7417 - loss: 0.1093 - val_accuracy: 0.7447 - val_loss: 0.1126
Epoch 6/20
114/114 140s 1s/step - accuracy: 0.7617 - loss: 0.1016 - val_accuracy: 0.7599 - val_loss: 0.1066
Epoch 7/20
114/114 131s 1s/step - accuracy: 0.7798 - loss: 0.0950 - val_accuracy: 0.7700 - val_loss: 0.1016
Epoch 8/20
114/114 125s 1s/step - accuracy: 0.7924 - loss: 0.0881 - val_accuracy: 0.7867 - val_loss: 0.0971
Epoch 9/20
114/114 140s 1s/step - accuracy: 0.8025 - loss: 0.0844 - val_accuracy: 0.7893 - val_loss: 0.0946
Epoch 10/20
114/114 127s 1s/step - accuracy: 0.8118 - loss: 0.0815 - val_accuracy: 0.7959 - val_loss: 0.0922
Epoch 11/20
114/114 125s 1s/step - accuracy: 0.8184 - loss: 0.0779 - val_accuracy: 0.7986 - val_loss: 0.0909
Epoch 12/20
114/114 142s 1s/step - accuracy: 0.8249 - loss: 0.0750 - val_accuracy: 0.8027 - val_loss: 0.0895
Epoch 13/20
114/114 140s 1s/step - accuracy: 0.8323 - loss: 0.0725 - val_accuracy: 0.8061 - val_loss: 0.0872
Epoch 14/20
114/114 127s 1s/step - accuracy: 0.8381 - loss: 0.0697 - val_accuracy: 0.8083 - val_loss: 0.0864
Epoch 15/20
114/114 142s 1s/step - accuracy: 0.8410 - loss: 0.0685 - val_accuracy: 0.8161 - val_loss: 0.0833
Epoch 16/20
114/114 132s 1s/step - accuracy: 0.8463 - loss: 0.0669 - val_accuracy: 0.8136 - val_loss: 0.0844
Epoch 17/20
114/114 130s 1s/step - accuracy: 0.8509 - loss: 0.0642 - val_accuracy: 0.8157 - val_loss: 0.0827
Epoch 18/20
114/114 143s 1s/step - accuracy: 0.8547 - loss: 0.0631 - val_accuracy: 0.8194 - val_loss: 0.0812
Epoch 19/20
114/114 134s 1s/step - accuracy: 0.8578 - loss: 0.0612 - val_accuracy: 0.8143 - val_loss: 0.0840
Epoch 20/20
114/114 132s 1s/step - accuracy: 0.8611 - loss: 0.0599 - val_accuracy: 0.8211 - val_loss: 0.0801
```

```
In [9]: import matplotlib.pyplot as plt
```

```
In [10]: plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label="Train Loss")
plt.plot(history.history['val_loss'], label="Val Loss")
plt.title("Loss (LSTM)")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label="Train Acc")
plt.plot(history.history['val_accuracy'], label="Val Acc")
plt.title("Accuracy (LSTM)")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



```
In [11]: model.save('pos_lstm_model.keras')
```

```
In [12]: test_loss, test_acc = model.evaluate(X_test, y_test, sample_weight=np.where(np.argmax(y_test, axis=-1) != 0, 1.0, 0.0))
print(f"\nTest Loss: {test_loss:.4f} | Test Accuracy: {test_acc:.4f}")
```

23/23 2s 98ms/step - accuracy: 0.8251 - loss: 0.0716

Test Loss: 0.0728 | Test Accuracy: 0.8246

```
In [13]: import numpy as np
y_pred = model.predict(X_test)
y_pred_ids = np.argmax(y_pred, axis=-1)
y_true_ids = np.argmax(y_test, axis=-1)

# Reverse Label index
idx2label = {i: l for l, i in dataset.label2id.items()}

# Show sample predictions
for i in range(3):
    print(f"\n--- Sentence {i+1} ---")
    for j in range(seq_len):
        word_vec = X_test[i, j]
        if np.all(word_vec == 0): continue # padding

        pred_label = idx2label.get(y_pred_ids[i, j]-1,
        true_label = idx2label.get(y_true_ids[i, j]-1,
        print(f'{j:2}: Pred: {pred_label:6} | True: {true_label:6}')
```

23/23 ————— 7s 203ms/step

--- Sentence 1 ---		
0: Pred: INTJ	True:	INTJ
1: Pred: PART	True:	PART
2: Pred: SCONJ	True:	SCONJ
3: Pred: INTJ	True:	SCONJ
4: Pred: ADP	True:	ADP
5: Pred: PART	True:	PART
6: Pred: ADV	True:	ADV
7: Pred: SCONJ	True:	SCONJ
8: Pred: ADJ	True:	ADJ
9: Pred: INTJ	True:	INTJ
10: Pred: PROPN	True:	PROPN

```
-- Sentence 2 --
0: Pred: ADJ | True: ADJ
1: Pred: SCONJ | True: INTJ
2: Pred: PART | True: PART
3: Pred: INTJ | True: SCONJ
4: Pred: CCONJ | True: CCONJ
5: Pred: INTJ | True: INTJ
6: Pred: ADJ | True: ADJ
7: Pred: INTJ | True: INTJ
8: Pred: ADJ | True: <PAD>
9: Pred: ADJ | True: ADJ
10: Pred: SCONJ | True: INTJ
11: Pred: AUX | True: AUX
12: Pred: CCONJ | True: CCONJ
13: Pred: INTJ | True: INTJ
14: Pred: PART | True: PART
15: Pred: PART | True: PART
16: Pred: SCONJ | True: SCONJ
17: Pred: INTJ | True: ADJ
18: Pred: CCN | True: PART
19: Pred: PROPN | True: PROPN
```

--- Sentence 3 ---		
0:	Pred: AUX	True: ADV
1:	Pred: CCONJ	True: CCONJ
2:	Pred: INTJ	True: INTJ
3:	Pred: ADJ	True: ADJ
4:	Pred: NOUN	True: NOUN
5:	Pred: INTJ	True: INTJ
6:	Pred: PROPN	True: PROPN
7:	Pred: INTJ	True: ADP
8:	Pred: SCONJ	True: INTJ
9:	Pred: AUX	True: AUX
10:	Pred: SCONJ	True: SCONJ
11:	Pred: ADJ	True: ADJ
12:	Pred: INTJ	True: INTJ
13:	Pred: ADJ	True: ADJ
14:	Pred: CCONJ	True: CCONJ
15:	Pred: INTJ	True: INTJ
16:	Pred: ADJ	True: ADJ
17:	Pred: CCONJ	True: CCONJ
18:	Pred: NOUN	True: NOUN
19:	Pred: INTJ	True: INTJ
20:	Pred: ADJ	True: ADJ
21:	Pred: INTJ	True: INTJ
22:	Pred: PROPN	True: PROPN

```
In [14]: np.argmax(y_test, axis=-1)[0]
```

```
In [15]: y_true_flat = y_true_ids.reshape((-1,))  
y_pred_flat = y_pred_ids.reshape((-1,))  
mask = (y_true_flat != 0)  
y_true_flat = y_true_flat[mask]  
y_pred_flat = y_pred_flat[mask]  
  
np.max(y_true_flat), np.min(y_true_flat), np.max(y_pred_flat), np.min(y_pred_flat)  
np.sum(y_true_flat==7), np.sum(y_pred_flat==7)
```

```
Out[15]: (np.int64(6), np.int64(0))
```

```
In [16]: from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
from matplotlib.colors import LogNorm

# Compute confusion matrix
labels = dataset.labels
label_ids = list(range(len(labels)))

print("Classification Report:")
print(classification_report(y_true_flat, y_pred_flat, labels=label_ids, target_names=labels))

cm = confusion_matrix(y_true_flat, y_pred_flat, labels=label_ids)

# Plot with log scale
plt.figure(figsize=(12, 10))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', norm=LogNorm(vmin=1, vmax=cm.max()),
            xticklabels=labels, yticklabels=labels, cbar_kws={'label': 'Log-scaled count'})
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Confusion Matrix (LSTM)")
plt.tight_layout()
plt.show()
```

Classification Report:				
	precision	recall	f1-score	support
<PAD>	0.00	0.00	0.00	0
ADJ	0.65	0.50	0.57	1172
ADP	0.95	0.98	0.96	2333
ADV	0.67	0.64	0.65	650
AUX	0.89	0.90	0.89	618
CCONJ	0.95	0.96	0.95	471
DET	0.84	0.81	0.82	898
INTJ	0.00	0.00	0.00	6
NOUN	0.74	0.85	0.79	4042
NUM	0.86	0.75	0.80	456
PART	0.94	0.90	0.92	358
PRON	0.81	0.74	0.77	862
PROPN	0.72	0.65	0.68	455
PUNCT	0.97	1.00	0.98	2083
SCONJ	0.73	0.77	0.75	154
VERB	0.79	0.70	0.74	1749
X	1.00	0.18	0.30	17
accuracy			0.82	16324
macro avg	0.74	0.67	0.68	16324
weighted avg	0.82	0.82	0.82	16324

```
d:\anu1m\sem2\NLP\soft\PoSTagging\.venv\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
d:\anu1m\sem2\NLP\soft\PoSTagging\.venv\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels with no true samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
d:\anu1m\sem2\NLP\soft\PoSTagging\.venv\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no true nor predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
d:\anu1m\sem2\NLP\soft\PoSTagging\.venv\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
d:\anu1m\sem2\NLP\soft\PoSTagging\.venv\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels with no true samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
d:\anu1m\sem2\NLP\soft\PoSTagging\.venv\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no true nor predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
d:\anu1m\sem2\NLP\soft\PoSTagging\.venv\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
d:\anu1m\sem2\NLP\soft\PoSTagging\.venv\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels with no true samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
d:\anu1m\sem2\NLP\soft\PoSTagging\.venv\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no true nor predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Confusion Matrix (LSTM)

